

—

THE AGENT

Your 24/7 analyst for retail

& consumer intelligence.

Goal: make you the smartest person on retail × AI in the room, any day, any time. Rachel is a Claude-powered research analyst. She reads investor filings the night they publish, scores AI adoption on a two-axis public rubric across a 50+ brand watchlist, and answers questions with a citation, or refuses to answer at all.

1

THE GROUNDED ANALYST

Answers with a citation, or not at all

- WHAT
- Ask anything on the watchlist. Hybrid retrieval over the ingested filings. Every claim carries a source and an evidence tier — confirmed, estimated, or absent.
- IMPACT
- Briefing-grade answers you can repeat in a boardroom. An ungrounded claim is treated as a bug, not a style problem: the refusal contract is enforced in the prompt and *tested by an eval harness on every release*.
- CONTEXT
- > No source, no claim · the refusal contract

> 0.98 gate · grounded-and-cited eval

> 3 retrievals max per question

HAIKU → SONNET → DEEP

MODEL ROUTING

2

THE SCORING FRAMEWORK

Two axes, four quadrants, every field sourced

- WHAT
- What the board says (AI rhetoric, 1–5) scored separately from what they actually shipped (confirmed production tools, 0–5+), by a mechanical rubric on public filings.
- IMPACT
- Silent builders surface — the companies doing real work and saying nothing, invisible to every single-score ranking. Narrative and production *never get conflated*.
- CONTEXT
- > METHODOLOGY.md · reproducible anchors

> CORRECTIONS.md · right of reply

> Every score dated to its filing

NEXT.JS · VERCEL

UPSTASH REDIS + VECTOR

3

THE INGESTION MACHINE

Filings read the night they publish

- WHAT
- A nightly job checks each brand's investor page against the earnings calendar, and on a new filing extracts 27 structured fields. Hard numbers apply directly; judgment calls queue for a human.
- IMPACT
- An earnings season reads itself. You review decisions, *not documents* — and walk into the next meeting already current.
- PART OF A FAMILY
- > Josh · career transition

> Holy · health & energy

> Cody · code auditor

VERCEL CRON

BLOB STORAGE

HUMAN-IN-THE-LOOP

—

THE LOOP

How Rachel reads an earnings season.

Filings land at midnight. By morning every affected profile is re-scored against the public rubric, with the human reviewing only the judgment calls. Ask a question and the answer cites the document it came from.

What Rachel reads

EXTERNAL

- Investor relations pages

WEB

Every brand's own IR site · checked nightly against the calendar
- Vendor case studies

WEB

Microsoft · Google · Anthropic · OpenAI · AWS naming their customers
- Retail & AI newsletters

WEB

Retail Dive · Modern Retail · BoF · Grocery Gazette · see config/
- Earnings calendar

API

Who reports when, so nothing is read late

INTERNAL

- One file per brand

CLI

Everything learned about that company, with the source behind each fact
- The searchable library

API

Every filing chunked and indexed, so a question finds its paragraph
- Daily pulse

API

Fresh signals from the newsletters, deduplicated, nothing over a week old
- Waiting for you

CLI

The judgment calls Rachel will not make alone

NIGHTLY PIPELINE

Check → extract → score → cite

Five steps, cheapest model that can do each job. Extraction runs on the Batches API at half token cost while you sleep.

- 01

Check

● TRIGGER

Nightly sweep of every brand's investor page, driven by the earnings calendar. New filing found and stored. No model calls, so this costs nothing.
- 02

Extract

● AGENTIC

Filing → 27 fields. Revenue, operating profit, headcount, store count, named AI tools, vendor partners, the CEO's exact words on AI. Each tagged confirmed, estimated, or absent.
- 03

Review

● TRIGGER

Hard numbers apply straight through. The judgment calls — is that tool really live? is this a moat or a cost story? — queue for a human.
- 04

Score

● TRIGGER

The rubric recomputes both axes and the quadrant. Every score stamped with the filing it came from.
- 05

Publish

● AGENTIC

Pulse items and answers, each with citations attached. No source in the library, no claim made.

MODEL ROUTING

Extraction & lookups → **Haiku** (Batches API, ~50% tokens). Reasoning, synthesis & judgment → **Sonnet**. Ranking the whole sector, or comparing many brands across quarters → **deep tier (Opus-class)**, opt-in via **DEEP_ANALYSIS_MODEL**. Eval gate: **0.98 grounded-and-cited**.

R

RACHEL

The brand file

27

FIELDS

·

< 90

SEC

One screen per brand, rebuilt every time they file. Each field tagged confirmed, estimated or absent — a blank cell and a researched absence must never look the same.

- 1

AI Impression Score

Both axes and the quadrant: Performer, Narrative-led, Silent builder, Not yet visible.
- 2

Investor snapshot

Revenue, operating profit, headcount, HQ and AI Act exposure, next reporting date.
- 3

AI perception

Named production tools, quantified ROI, consumer-facing AI, vendor partners, who presented it.
- 4

What the CEO said

The exact sentence the rhetoric score is anchored to, with the call and the date.
- 5

Recent sources

Every filing and case study behind a claim, dated and linked.
- 6

Ask Rachel

NEW

Scoped Q&A on this brand. Citation or refusal, nothing in between.